



Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Query Optimisation

Assignment Project Exam Help

- In practice, query optimisers incorporate elements of the following three optimisation approaches:

• <https://eduassistpro.github.io>

- **Rule-based query optimisation**

Use heuristic rules to transform a relational equivalent one with a possibly lower cost.

- **Cost-based query optimisation**

Use a cost model to estimate the costs of plans, and then select the most cost-effective plan.



Semantic Query Optimisation

Assignment Project Exam Help

- Can we use semantic information stored in a database (such as integrity constraints) to optimise queries?



- R <https://eduassistpro.github.io>



- entity integrity constraints
- referential integrity constraints
- domain constraints



- ...
- user-defined integrity constraints

- **Key idea:** Integrity constraints may **not only be utilized to enforce consistency** of a database, but may **also optimise user queries**.



Semantic Query Optimisation

Assignment Project Exam Help

• E

<https://eduassistpro.github.io>

Query: `SELECT DISTINCT ssn FROM Empl`

Add WeChat `edu_assist_pr`

- We can avoid extra costs for duplicate elimi
constraint tells us that tuples in the result will be unique.

Semantic Query Optimisation

Assignment Project Exam Help

- Example 2:

<https://eduassistpro.github.io>

```
FROM Employee
```

```
WHERE salary > 300000
```

Add WeChat edu_assist_pro

- We do not need to execute a query if the existing constraint tells us that the result will be empty.



Semantic Query Optimisation

Assignment Project Exam Help

• Example 3:

<https://eduassistpro.github.io>

```
FROM Works_on INNER JOIN Project
```

```
on Works_on.pno=Proje
```

• We can reduce the number of joins by executing
since both queries always return the same r

```
SELECT DISTINCT ssn
```

```
FROM Works_on ;
```



Rule-based Query Optimisation

Assignment Project Exam Help

- A rule-based optimisation transforms the RA expression by using a set of heuristic rules that typically improve the execution performance.

- **K**
w

<https://eduassistpro.github.io>

Apply as early as possible to reduce the number of tuples;

- **Push-down projection:**

Apply as early as possible to reduce the num

- **Re-ordering joins:**

Apply restrictive joins first to reduce the size of the result.

- But we must ensure that the resulting query tree gives the same result as the original query tree, i.e., **the equivalence of RA expressions.**



Heuristic Rules

Assignment Project Exam Help

Staff(staffNo, name, salary, position, branchNo)
Branch(branchNo, name, street, suburb, city)

- T by <https://eduassistpro.github.io> s, utilized

$$(1) \sigma_{\varphi}(\sigma_{\psi}(R)) \equiv \sigma_{\varphi \wedge \psi}(R);$$

$$\sigma_{branchNo='1' \wedge (\sigma_{salary > 60000}(Staff))} =$$

$$(2) \pi_X(\pi_Y(R)) \equiv \pi_X(R) \text{ if } X \subseteq Y$$

$$\pi_{salary}(\pi_{branchNo, salary}(Staff)) = \pi_{salary}$$

$$(3) \sigma_{\varphi}(R_1 \times R_2) \equiv R_1 \bowtie_{\varphi} R_2$$

$$\sigma_{Staff.branchNo=Branch.branchNo}(Staff \times Branch) =$$

$$(Staff) \bowtie_{Staff.branchNo=Branch.branchNo} (Branch)$$



Heuristic Rules

Assignment Project Exam Help

Staff (sId, lName, fName, salary, position, branchNo)
Branch(branchNo, name, street, suburb, city)

<https://eduassistpro.github.io>

$(Staff) \bowtie_{Staff.branchNo=Branch.br}$

(5) $\sigma_{\varphi}(R_1 \bowtie R_2) \equiv \sigma_{\varphi}(R_1) \bowtie R_2$, if

$\sigma_{salary > 60000}(Staff \bowtie Branch) = \sigma_{sal}$

(6) $\sigma_{\varphi_1 \wedge \varphi_2}(R_1 \bowtie R_2) \equiv \sigma_{\varphi_1}(R_1) \bowtie \sigma_{\varphi_2} R_2$ if φ_1 contains only attributes in R_1 and φ_2 contains only attributes in R_2 .

$\sigma_{salary > 60000 \wedge city = 'Canberra'}(Staff \bowtie Branch) =$

$(\sigma_{salary > 60000}(Staff)) \bowtie (\sigma_{city = 'Canberra'}(Branch))$



Heuristic Rules

Assignment Project Exam Help

Staff(staffNo, lname, fname, salary, position, branchNo)
Branch(branchNo, name, street, suburb, city)

<https://eduassistpro.github.io>

$\pi_{branchNo, position, city}(Staff \bowtie Branch)$

Add WeChat: edu_assist_pro

(3) If the join condition contains attributes not in both R_1 and R_2 , and ones in both R_i and X

$\pi_{position, city}(Staff \bowtie Branch) =$

$\pi_{position, city}(\pi_{branchNo, position}(Staff) \bowtie (\pi_{branchNo, city}(Branch)))$



Push-down Selection – Example

Assignment Project Exam Help

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

DI

M

<https://eduassistpro.github.io>

- Query:** List the first and last names of the direct a movie that has won an 'Oscar' movie award

$\pi_{first_name, last_name}(\sigma_{award_name = 'Oscar'}((\sigma_{PERSON} \bowtie MOVIE_AWARD)))$

- Question:** Can we apply the following rule to optimise the query?

$\sigma_{\varphi}(R_1 \bowtie R_2) \equiv \sigma_{\varphi}(R_1) \bowtie R_2$, if φ contains only attributes in R_1

Add WeChat edu_assist_pr



Push-down Selection – Example

Assignment Project Exam Help

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

DI

M

<https://eduassistpro.github.io>

- Query:** List the first and last names of the direct a movie that has won an 'Oscar' movie award

$\pi_{first_name, last_name}(\sigma_{award_name='Oscar'}((PERSON \bowtie DIRECTOR) \bowtie MOVIE_AWARD))$

- We would have:

$\pi_{first_name, last_name}((PERSON \bowtie DIRECTOR) \bowtie \sigma_{award_name='Oscar'}(MOVIE_AWARD))$



Push-down Projection – Example

Assignment Project Exam Help

- Given the relation schemas.

PERSON(id, first_name, last_name, year_born)

DI

M

(id, first_name, last_name, year_born, director_id)

<https://eduassistpro.github.io>

- Q

a movie that has won an 'Oscar' movie award

$\pi_{first_name, last_name}((PERSON \bowtie DIRECTOR \bowtie AWARD))$

- Question:** Can we apply the following rule to optim

$$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)),$$

where X_i contains attributes in both in R_1 and R_2 , and ones in both R_i and X



Push-down Projection – Example

Assignment Project Exam Help

- Given the relation schemas.

PERSON(id, first_name, last_name, year_born)

DI

M

(id, first_name, last_name, year_born, title, production_year, award_name)

<https://eduassistpro.github.io>

- Q a movie that has won an 'Oscar' movie award

$\pi_{first_name, last_name}((PERSON \bowtie DIRECTOR \bowtie MOVIE_AWARD))$

- we would have:

$\pi_{first_name, last_name}(\pi_{first_name, last_name, title, production_year}(PERSON \bowtie DIRECTOR) \bowtie \pi_{title, production_year}(\sigma_{award_name='Oscar'}(MOVIE_AWARD)))$



A Common Query Pattern (Be Careful)

Assignment Project Exam Help

- A common query pattern is **join select-project** involving three steps:
 - (1) **join** all the relevant relations,

- <https://eduassistpro.github.io>

$$\pi_{A_1, \dots, A_n}(\sigma_{\varphi}(R_1 \times \dots \times R_k))$$

or as an equivalent SQL statement

Add WeChat edu_assist_pro

```
SELECT DISTINCT A1, ..., An FROM R1, ..., Rk WHERE  $\varphi$ ;
```

- Queries falling into this pattern can be **very inefficient**, which may yield huge intermediate result for the joined relations.



A Common Query Pattern (Be Careful)

push-down selection and push-down projection.

Assignment Project Exam Help

$$\pi_{A_1, \dots, A_n}(\sigma_{\varphi}(R_1 \times \dots \times R_k)),$$

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Re-ordering Joins - Example

Assignment Project Exam Help

- Given the relation schemas:
`PERSON(id, first_name, last_name, year_born)`

Suppose that it has **10000 tuples**.

`DI`

<https://eduassistpro.github.io>

Suppose that it has **100 tuples**.

`MOVIE_AWARD(title, production_year, award)`

Suppose that it has **1000 tuples**.

Add WeChat edu_assist_pro

- Example:** Consider the following two RA queries. Which one is better?

- `PERSON ⋈ MOVIE_AWARD ⋈ DIRECTOR`
- `PERSON ⋈ DIRECTOR ⋈ MOVIE_AWARD`

Cost-based Query Optimisation

Assignment Project Exam Help

- A query optimizer

<https://eduassistpro.github.io>

- It estimates and compares the costs of executing a query using different execution strategies and chooses one with the lowest cost.

Add WeChat edu_assist_pro

- The query optimiser needs to **limit the number of plans** to be considered for improving efficiency.



Summary

Assignment Project Exam Help

- In general, there are many ways of executing a query in a database

- T
b

- B
sh

<https://eduassistpro.github.io>

- Nonetheless, SQL is not a suitable query language optimised automatically

Add WeChat edu_assist_pro

- Instead, SQL queries are **transformed into queries** and optimised subsequently.

- A major advantage of relational algebra is to **make alternative forms of a query easy to explore.**