



Week 8 Workshop – Query Processing and Optimisation

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Housekeeping

Assignment Project Exam Help

- Assignment 2 (Database Theory) for both COMP2400/6240 students

- The submission deadline is 23:59, Oct 12, 2021.
-

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Housekeeping

Assignment Project Exam Help

- Assignment 2 (Database Theory) for both COMP2400/6240 students

- The submission deadline is 23:59, Oct 12, 2021.
-

② All f
sa <https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Housekeeping

Assignment Project Exam Help

- Assignment 2 (Database Theory) for both COMP2400/6240 students

- The submission deadline is 23:59, Oct 12, 2021.
-

- All f
sa

- Lab 8 is optional (no associated with any assessme

- We will open a separate sign-up page on 'Mat'
 - All the optional labs will be scheduled from O
 - Three options are available
 - (1) Database Programming with Java
 - (2) Database Programming with Python
 - (3) Database Exercises on IMDB

Add WeChat edu_assist_pr



Query Processing – Example

SELECT name FROM Person WHERE age < 21

High-level language
(SQL)



uage
ebra)

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



name
Rickon
Bran

Query result



From SQL to RA Expressions

Assignment Project Exam Help

```
Students(matNr, firstName, lastName, email)
Exams(matNr, crsNr, result, semester)
Courses(crsNr, title, unit)
```

<https://eduassistpro.github.io>

```
EXAMS.crsNr=COURSES.crsNr AND result 1.3;
```

Add WeChat edu_assist_pro



From SQL to RA Expressions

Assignment Project Exam Help

Students(matNr, firstName, lastName, email)
Exams(matNr, crsNr, result, semester)
Courses(crsNr, title, unit)

<https://eduassistpro.github.io>

EXAMS.crsNr=COURSES.crsNr AND result 1.3;

- RA Expressions:

- 1 $\pi_{lastName, result, title}(\sigma_{result \leq 1.3}((Student \bowtie_{Exams.crsNr=Courses.crsNr} Courses)))$
- 2 $\pi_{lastName, result, title}(\sigma_{result \leq 1.3}(\sigma_{ExAMS.crsNr=Courses.crsNr}(\sigma_{Students.matNr=Exams.matNr}(Students \times Exams \times Courses))))$
- 3 $\pi_{lastName, result, title}((Students \bowtie_{Students.matNr=Exams.matNr} (\sigma_{result \leq 1.3}(Exams))) \bowtie_{Exams.crsNr=Courses.crsNr} Courses)$
- 4 ...



From RA Expressions to Query Trees

- Each RA expression can be represented as a **query tree**:
 - leaf nodes represent the input relations;
 - **internal nodes** represent the intermediate result;
 - ...

• $E \pi_{la} \bowtie$ <https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Query Tree Example

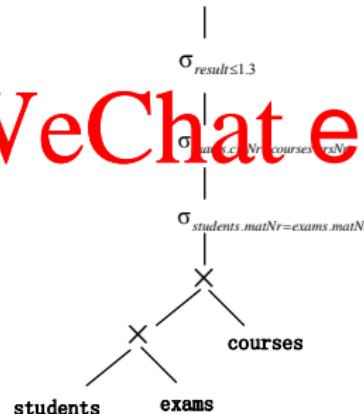
- For each query tree, computation proceeds bottom-up:
 - child nodes must be executed before their parent nodes.
 - but there can exist multiple methods of executing sibling nodes.

- E

π_{la}

$(\sigma$

<https://eduassistpro.github.io>



Add WeChat edu_assist_pro



Equivalent Query Trees (Query Optimisation)

Assignment Project Exam Help

$$\pi_{\text{lastName}, \text{result}, \text{title}}$$

<https://eduassistpro.github.io>

Add WeChat `edu_assist_pro`





Execution Plan

(Slide 8-27 will not be assessed in our course)

Assignment Project Exam Help

- A **query execution plan** consists of an (extended) query tree with additional annotation at each node indicating:

{
{

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Execution Plan

- **Materialized:** The intermediate result of an operator may be saved in a temporary table for processing by the next operator.
- **Pipelined:** the intermediate result of an operator is directly sent to another operator without creating a temporary table (also called **on-the-fly**).

<https://eduassistpro.github.io>

Add WeChat `edu_assist_pro`

Note: Pipelined evaluation may have significant saving on I/O cost, while materialized evaluation can avoid repeated computations.



Execution Plan

Assignment Project Exam Help

- Question: Which execution plan is "optimal" in terms of processing efficiency?

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Execution Plan

Assignment Project Exam Help

- Question: Which execution plan is "optimal" in terms of processing efficiency?

- Th

(F

<https://eduassistpro.github.io>

- R

consistently find a plan that is good.

Add WeChat edu_assist_pro



Execution Plan

Assignment Project Exam Help

- Question: Which execution plan is "optimal" in terms of processing efficiency?

• Th

(F

<https://eduassistpro.github.io>

• R

consistently find a plan that is good.

- The performance of different execution plans can differ considerably (e.g., seconds vs. hours vs. days):

- different but equivalent **RA expressions**;
- different algorithms for **each RA operator**.



Execution Plan

- Basic ideas of algorithms used for RA operators

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Execution Plan

Assignment Project Exam Help

- Basic ideas of algorithms used for RA operators
 - **Selection:** If there is no index we have to scan the table. Otherwise, we scan the indexes to retrieve matching tuples and apply remaining

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Execution Plan

Assignment Project Exam Help

- Basic ideas of algorithms used for RA operators
 - Selection: If there is no index we have to scan the table. Otherwise, we scan the indexes to retrieve matching tuples and apply remaining

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Execution Plan

Assignment Project Exam Help

- Basic ideas of algorithms used for RA operators
 - **Selection:** If there is no index we have to scan the table. Otherwise, we scan the indexes to retrieve matching tuples and apply remaining

<https://eduassistpro.github.io>

- **Join:** We may use *nested loops join*, *block nested loops join*, etc.

Add WeChat edu_assist_pro



Execution Plan

Assignment Project Exam Help

- Basic ideas of algorithms used for RA operators
 - **Selection:** If there is no index we have to scan the table. Otherwise, we scan the indexes to retrieve matching tuples and apply remaining

<https://eduassistpro.github.io>

- **Join:** We may use *nested loops join*

block nested loops join, etc.

s,

Add WeChat edu_assist_pro

- **Group by** and **order by** are typically I



Execution Plan

Assignment Project Exam Help

- Basic ideas of algorithms used for RA operators
 - **Selection:** If there is no index we have to scan the table. Otherwise, we scan the indexes to retrieve matching tuples and apply remaining

<https://eduassistpro.github.io>

- **Join:** We may use *nested loops join*, *block nested loops join*, etc.
- **Group by** and **order by** are typically implemented using temporary tables.
- **Aggregation operators** use temporary counters in main memory when retrieving tuples.

Add WeChat edu_assist_pro



Execution Plan

- Basic ideas of algorithms used for RA operators
 - **Selection:** If there is no index, we have to scan the table. Otherwise, we scan the indexes to retrieve matching tuples and apply remaining

<https://eduassistpro.github.io>

- **Join:** We may use *nested loops join*, *block nested loops join*, etc.
- **Group by** and **order by** are typically implemented using temporary files.
- **Aggregation operators** use temporary counters in main memory when retrieving tuples.
- **Set operators** can use the same approach as projection to eliminate duplicates.

Add WeChat `edu_assist_pro`



Estimating Query Costs - Example

Assignment Project Exam Help

- Which movies got a non-US award for one of its actors playing an ‘agent’?

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Estimating Query Costs - Example

Assignment Project Exam Help

- Which movies got a non-US award for one of its actors playing an 'agent'?

$$\pi_{\text{title}, \text{production_year}}(\sigma_{\text{role_description}='agent'}(\text{ROLE} \bowtie \text{ACTOR_AWARD} \bowtie (\text{AWARD} \sigma_{\text{award_country}='USA'} (\text{AWARD}))))$$

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Estimating Query Costs - Example

Assignment Project Exam Help

- Which movies got a non-US award for one of its actors playing an 'agent'?

$$\pi_{\text{title}, \text{production_year}}(\sigma_{\text{role_description}='agent'}(\text{ROLE} \bowtie \text{ACTOR_AWARD} \bowtie (\text{AWARD} \sigma_{\text{award_country}='USA'} (\text{AWARD}))))$$

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Size of Relations

Assignment Project Exam Help

- How to determine the size of a relation r over $F(A_1, \dots, A_n)$?

- Let n denote the average number of tuples in r , and ℓ_j the the average

<https://eduassistpro.github.io>

1				
...
...
...
...

Add WeChatedu_assist_pro



Size of Relations

Assignment Project Exam Help

- How to determine the size of a relation r over $F(A_1, \dots, A_k)$?

- Let n denote the average number of tuples in r , and ℓ_j the the average

<https://eduassistpro.github.io>

1				
...
...
...
...

Add WeChatedu_assist_pro

- Then, $n \cdot \sum_{j=1}^k \ell_j$ is the size of the relation r .
- We use this formula to assign sizes to leaf nodes in the query tree.

Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

AWARD: Award name: var char(30), Institution: var char(50), Award country:

varchar(20))

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

AWARD_Award_name: varchai(30), Institution:vachar(50), Award_country:

varchar(20))

• <https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

• AWARD(Award_name: varchar(30),Institution:varchar(50),Award_country:
varchar(20))

• <https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

- AWARD(Award_name: varchar(30),Institution:varchar(50),Award_country:
 varchar(20))

• <https://eduassistpro.github.io>

- Award_name: $8 \cdot 20 = 160$ bits (t)
- Institution: $8 \cdot 30 = 240$ bits (the)
- Award_country: $8 \cdot 10 = 80$ bits (c)

Add WeChat `edu_assist_pro`



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

- AWARD(Award_name: varchar(30),Institution:varchar(50),Award_country:
 varchar(20))

• <https://eduassistpro.github.io>

- Award_name: $8 \cdot 20 = 160$ bits (t)
 - Institution: $8 \cdot 30 = 240$ bits (the)
 - Award_country: $8 \cdot 10 = 80$ bits (c)
-
- Add WeChat `edu_assist_pro`
- The average size of a tuple is $160 + 80 + 240 = 480$ bits.



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

- AWARD(Award_name: varchar(30), Institution: varchar(50), Award_country:
 varchar(20))

<https://eduassistpro.github.io>

- Award_name: $8 \cdot 20 = 160$ bits (t)
- Institution: $8 \cdot 30 = 240$ bits (the)
- Award_country: $8 \cdot 10 = 80$ bits (c)

Add WeChat `edu_assist_pro`

- The average size of a tuple is $160 + 80 + 240 = 480$ bits.
- The average size of a relation is estimated to be $15 \cdot 480 = 7,200$ bits.



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

Role:IC:char(4),Title:varchar(40),Production year: number(4),

Role_description:varchar(100),Credits:varchar(40))

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

Role:IC:char(4),Title:varchar(40),Production year: number(4),

Role_description:varchar(100),Credits:varchar(40))

•

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

Role:IC:char(4),Title:varchar(40),Production year: number(4),
Role_description:varchar(100),Credits:varchar(40))

•

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

- Role: ID:char(4), Title:varchar(40), Production year: number(4),

Role_description:varchar(100), Credits:varchar(40))

• <https://eduassistpro.github.io>

- Title: $8 \cdot 25 = 200$ bits (the mean length is 25).
- Production year: 15 bits (as the domain has 100 years).
- Role_description: $8 \cdot 50 = 400$ bits.
- Credits: $8 \cdot 20 = 160$ bits (the mean length is 20).

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

- Role: ID:char(4), Title:varchar(40), Production year: number(4),

Role_description:varchar(100), Credits:varchar(40))

• <https://eduassistpro.github.io>

• Title: $8 \cdot 25 = 200$ bits (the mean length is 25).

• Production year: 13 bits (as the domain has 13 years).

• Role_description: $8 \cdot 50 = 400$ bits.

- Credits: $8 \cdot 20 = 160$ bits (the mean length is 20).
- The average size of a tuple is $64 + 200 + 13 + 400 + 160 = 837$ bits

Add WeChat `edu_assist_pro`



Estimating Query Costs - Example (Relation Sizes)

Assignment Project Exam Help

Role:ID:char(4),Title:varchar(40),Production year:number(4),

Role_description:varchar(100),Credits:varchar(40))

• <https://eduassistpro.github.io>

- Title: $8 \cdot 25 = 200$ bits (the mean length is 25).
 - Production year: 13 bits (as the domain has 50 years).
 - Role_description: $8 \cdot 50 = 400$ bits.
 - Credits: $8 \cdot 20 = 160$ bits (the mean length is 20).
- The average size of a tuple is $64 + 200 + 13 + 400 + 160 = 837$ bits
 - The average size of a relation is to be $500 \cdot 837 = 418,500$ bits

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Relation Sizes)

- ACTOR_AWARD(Title:varchar(40),Production_year:number(4),

Role_description:varchar(100),Award_name:varchar(30),

Year_of_award:number(4),Category:varchar(100),Result:varchar(20))

<https://eduassistpro.github.io>

- Production_year: 13 bits (as before);
- Role_description: 400 bits (as before);
- Award_name: 100 bits (as before);
- Year_of_award: 13 bits (as the domain);
- Category: $8 \cdot 40 = 320$ bits (the mean length is 40);
- Result: $8 \cdot 7 = 56$ bits (the mean length is 7).

- The average size of a tuple is $200 + 13 + 400 + 160 + 13 + 320 + 56 = 1,162$ bits.
- The average size of a relation is $40 \cdot 1162 = 46,480$ bits.

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Query Tree)

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Size of Selection Node

Assignment Project Exam Help

- Selection π_{φ} is linear in the number n of tuples of the involved relation

•

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Size of Selection Node

Assignment Project Exam Help

- Selection σ_φ is linear in the number n of tuples of the involved relations

•

<https://eduassistpro.github.io>

- Let s be the size of its single relevant node.
- The size of a selection node σ_φ is

$$a_\varphi \cdot s ,$$

where a_φ is the average percentage of tuples satisfying φ .

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Selection)

Assignment Project Exam Help

- For selection $\text{award_country} = \text{'USA'}$ assume $a_{\varphi} = 0.4$ i.e., 40% of the movie awards from the USA). Hence, we have: $a_{\varphi} \cdot s = 0.4 \cdot 7,200 = 2,880$.

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Size of Difference Node

Assignment Project Exam Help

- Let s_1 and s_2 be the sizes of the two relevant nodes.

- A rel <https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Size of Difference Node

Assignment Project Exam Help

- Let s_1 and s_2 be the sizes of the two relevant nodes.

- A rel <https://eduassistpro.github.io>

- The size of a difference node is

Add WeChat $\frac{s_1}{s_1 + s_2} - p$

where $(1 - p)$ is the probability that tuples from s_1 does not occur in s_2 .



Estimating Query Costs - Example (Difference)

- Since 40% of the movie awards from the USA, the probability of an award to be a US award is $p = 0.4$. We have: $S_1 \cdot (1 - p) = 7,200 \cdot (1 - 0.4) = 4,320$.

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Size of Natural Join Node

Assignment Project Exam Help

- Let s_1 and s_2 be the sizes of the two relevant nodes, and r_1 and r_2 be the

si

$\underline{s_1}$

$\underline{s_2}$

ber of

tu

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Size of Natural Join Node

Assignment Project Exam Help

- Let s_1 and s_2 be the sizes of the two relevant nodes, and r_1 and r_2 be the

si

$\underline{s_1}$

$\underline{s_2}$

ber of

tu

<https://eduassistpro.github.io>

- Th

Add WeChat edu_assist_pro

$$\frac{s_1 \cdot s_2}{r_1 + r_2 - r} (1 - p)$$

where r is the size of a tuple over the **co**

matching probability (for any tuple of the first relevant node to match with any tuples in the second relevant relation). Note that $r_1 + r_2 - r$ is the size of a tuple after the natural join operation.



Estimating Query Costs - Example (Natural Join)

- For join with ACTOR_AWARD assume $p = 0.08$, i.e., 8% of the actor awards are non-US awards. By $\frac{s_1}{r_1} \cdot p \cdot \frac{s_2}{r_2} \cdot (r_1 + r_2 - r)$, we have:

4

1

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Natural Join)

Assignment Project Exam Help

- Assume $p = 0.002$. By $\frac{s_1}{r_1} \cdot p \cdot \frac{s_2}{r_2} \cdot (l_1 + l_2 - r)$, we have:

418,500

42,682

9,133.

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Estimating Query Costs - Example (Selection)

- For selection $\sigma_{\text{role_description} = \text{'agent'}}$, assume $a_{10} = 0.05$ (i.e., non-US awards for "agent" roles are 5%). Hence, we have: $a_{10} \cdot s = 0.05 \cdot 49133 = 2,457$.

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Size of Projection Node

Assignment Project Exam Help

- $\text{Projection}(\pi_{\{A_1, \dots, A_n\}})$:
- Project each tuple to the attributes in $\{A_1, \dots, A_n\}$
-

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Size of Projection Node

Assignment Project Exam Help

- Projection $\pi_{A_1, \dots, A_n} \{ \cdot \}$:
 - Project each tuple to the attributes in $\{A_1, \dots, A_n\}$
 -
- Let <https://eduassistpro.github.io>
- The size of a projection node $\pi_{A_1, \dots, A_n} \{ \cdot \}$ is

Add WeChat `edu_assist_pr`

$$(1 - p_i) \cdot s -$$

where r_i is the average size of a tuple over $\{A_1, \dots, A_n\}$, and p_i is the probability that two tuples coincide on A_1, \dots, A_n (i.e., the same values on all attributes A_1, \dots, A_n).



Estimating Query Costs - Example (Projection)

- For projection $\pi_{\text{title}, \text{production_year}}$, assume that there are 1% of duplicates, i.e., $p_i = 0.01$. By $(1 - p_i) \cdot s \cdot \frac{r_i}{r}$, we have $(1 - 0.01) \cdot 2,457 \frac{2,312,304}{1706} = 304$

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Query Optimisation

Assignment Project Exam Help

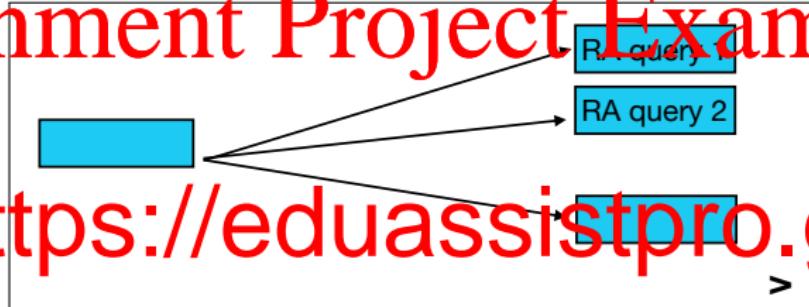
<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Relational Algebra \Rightarrow Query Optimisation

Assignment Project Exam Help



<https://eduassistpro.github.io>

>

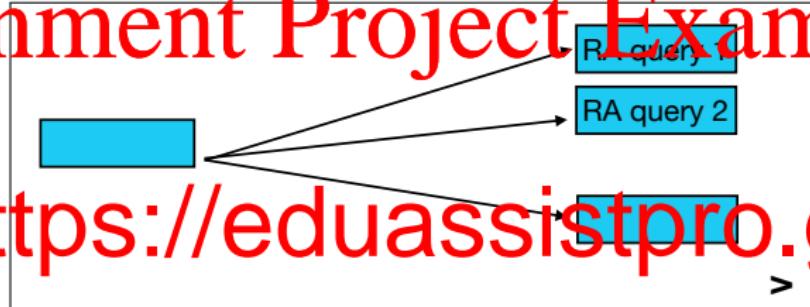
- Which RA query should be chosen for a given SQL query?

Add WeChat edu_assist_pro



Relational Algebra \Rightarrow Query Optimisation

Assignment Project Exam Help



<https://eduassistpro.github.io>

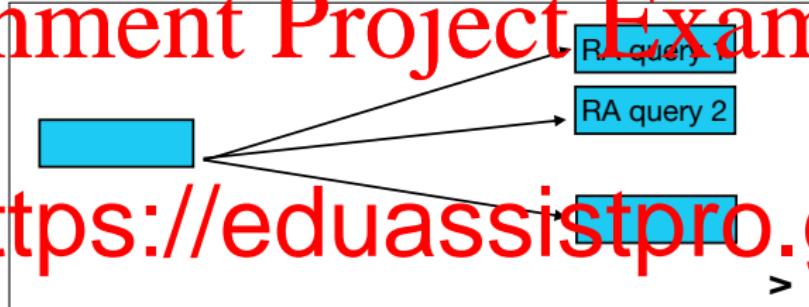
- Which RA query should be chosen for a given SQL query?
 - Who chooses?

Add WeChat **edu_assist_pro**



Relational Algebra \Rightarrow Query Optimisation

Assignment Project Exam Help



<https://eduassistpro.github.io>

- Which RA query should be chosen for a given SQL query?

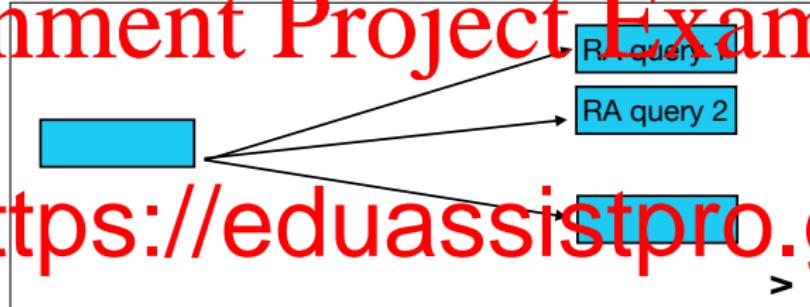
Add WeChat `edu_assist_pro`

- Who chooses? Query optimiser!



Relational Algebra \Rightarrow Query Optimisation

Assignment Project Exam Help



<https://eduassistpro.github.io>

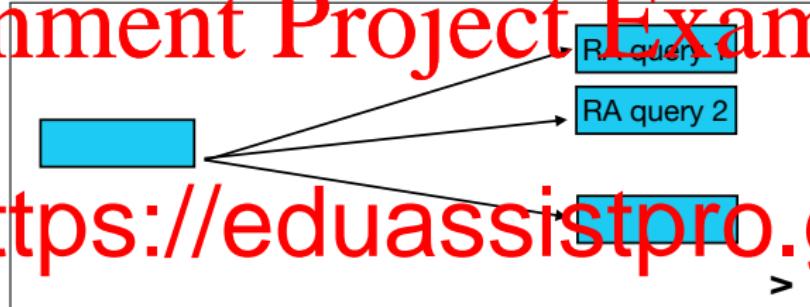
- Which RA query should be chosen for a given SQL query?
 - Who chooses? *Query optimiser!*
 - How to choose?

Add WeChat `edu_assist_pro`



Relational Algebra \Rightarrow Query Optimisation

Assignment Project Exam Help



<https://eduassistpro.github.io>

- Which RA query should be chosen for a given SQL query?

- Who chooses? **Query optimiser!**
 - How to choose?

- Semantic query optimisation
 - Rule-based optimisation
 - Cost-based optimisation

Add WeChat **edu_assist_pro**



Query Optimisation

Assignment Project Exam Help

- In practice, query optimisers incorporate elements of the following three optimisation approaches:

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Query Optimisation

Assignment Project Exam Help

- In practice, query optimisers incorporate elements of the following three optimisation approaches:

•

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Query Optimisation

Assignment Project Exam Help

- In practice, query optimisers incorporate elements of the following three optimisation approaches:



<https://eduassistpro.github.io>

- **Rule-based query optimisation**

Add WeChat edu_assist_pro

Use heuristic rules to transform a relational query into an equivalent one with a possibly lower cost.



Query Optimisation

Assignment Project Exam Help

- In practice, query optimisers incorporate elements of the following three optimisation approaches:



<https://eduassistpro.github.io>

- **Rule-based query optimisation**

Add WeChat edu_assist_pro

Use heuristic rules to transform a relational query into an equivalent one with a possibly lower cost.

- **Cost-based query optimisation**

Use a cost model to estimate the costs of plans, and then select the most cost-effective plan. This will not be assessed in our course.



Semantic Query Optimisation

Assignment Project Exam Help

- Example:

PERSON(id, first_name, last_name, year_born)

— — — — —)

<https://eduassistpro.github.io>

- List the ids of the writers who have written movies pro

Add WeChat edu_assist_pro



Semantic Query Optimisation

Assignment Project Exam Help

- Example:

PERSON(id, first_name, last_name, year_born)

— — — — —)

<https://eduassistpro.github.io>

- List the ids of the writers who have written movies produced in 2000

$\pi_{id} \sigma_{production_year=2000} (WRITER \bowtie PERSON)$

$\pi_{id} \sigma_{production_year=2000} (WRITER \bowtie PERSON)$

$\pi_{id} \sigma_{production_year=2000} (WRITER \bowtie MOVIE)$

Add WeChat edu_assist_pro



Semantic Query Optimisation

Assignment Project Exam Help

- Example:

$\text{PERSON}(\underline{id}, \text{first_name}, \text{last_name}, \text{year_born})$

<https://eduassistpro.github.io>

- List the ids of the writers who have written movies produced in 2000

$\pi_{id} \sigma_{\text{production_year}=2000} (\text{WRITER} \bowtie \text{PERSON})$

$\pi_{id} \sigma_{\text{production_year}=2000} (\text{WRITER} \bowtie \text{PERSON})$

$\pi_{id} \sigma_{\text{production_year}=2000} (\text{WRITER} \bowtie \text{MOVIE})$

- $\pi_{id} \sigma_{\text{production_year}=2000} \text{WRITER}$ ← the optimised RA

Add WeChat edu_assist_pro



Rule-based Query Optimisation

Assignment Project Exam Help

- A rule-based optimisation transforms the RA expression by using a set of heuristic rules that typically improve the execution performance.

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Query Optimisation

Assignment Project Exam Help

- A rule-based optimisation transforms the RA expression by using a set of heuristic rules that typically improve the execution performance.

• ^K
^w

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Query Optimisation

Assignment Project Exam Help

- A rule-based optimisation transforms the RA expression by using a set of heuristic rules that typically improve the execution performance.

• ^K
^w

<https://eduassistpro.github.io>

.

Apply as early as possible to reduce the num

Add WeChat edu_assist_pro



Rule-based Query Optimisation

Assignment Project Exam Help

- A rule-based optimisation transforms the RA expression by using a set of heuristic rules that typically improve the execution performance.

• K

w <https://eduassistpro.github.io>

.

Apply as early as possible to reduce the number of attributes.

Add WeChat edu_assist_pro

Apply as early as possible to reduce the number of attributes.



Rule-based Query Optimisation

Assignment Project Exam Help

- A rule-based optimisation transforms the RA expression by using a set of heuristic rules that typically improve the execution performance.

• K

w <https://eduassistpro.github.io>

.

Apply as early as possible to reduce the number of attributes.

Add WeChat edu_assist_pro

Apply as early as possible to reduce the number of attributes.



Rule-based Query Optimisation

Assignment Project Exam Help

- A rule-based optimisation transforms the RA expression by using a set of heuristic rules that typically improve the execution performance.

• K

w <https://eduassistpro.github.io>

.

Apply as early as possible to reduce the number of attributes.

Add WeChat edu_assist_pro

Apply as early as possible to reduce the number of attributes.

- But we must ensure that the resulting query tree gives the same result as the original query tree, i.e., **the equivalence of RA expressions**.



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? ↪ Merging RA operators.

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? ↪ Merging RA operators.

• <https://eduassistpro.github.io>

- $\sigma_\varphi(R_1 \times R_2) \equiv R_1 \bowtie_\varphi R_2$

- $\sigma_{\varphi_1}(R_1 \bowtie_{\varphi_2} R_2) = R_2 \bowtie_{\varphi_1 \circ \varphi_2} R_1$

• Add WeChat edu_assist_pro



Rule-based Optimisation

- Can they be executed in one go? ↗ Merging RA operators

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can they be executed in one go? \hookrightarrow Merging RA operators

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can they be executed in one go? \hookrightarrow Merging RA operators

Assignment Project Exam Help

$$\sigma_{\varphi}(\sigma_{\psi}(R)) \equiv \sigma_{\varphi \wedge \psi}(R);$$

$\sigma_{CourseNo}$

STUDY)

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can they be executed in one go? \hookrightarrow Merging RA operators

Assignment Project Exam Help

$$\sigma_{\varphi}(\sigma_{\psi}(R)) \equiv \sigma_{\varphi \wedge \psi}(R);$$

$\sigma_{CourseNo}$

UID	CourseNo	Hours
111		
222	COMP2400	115
333	STAT2001	120
111	BUSN2011	110
111	ECON2101	120
333	BUSN2011	130

UID	CourseNo	Hours
111	COMP2400	120
111	BUSN2011	
111	ECON2102	

STUDY		
UID	CourseNo	Hours
0		120

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can they be executed in one go? \rightarrow Merging RA operators

Assignment Project Exam Help

$$\sigma_{\varphi}(\sigma_{\psi}(R)) \equiv \sigma_{\varphi \wedge \psi}(R);$$

$\sigma_{CourseNo}$

UID	STUDY	
111		
222	COMP2400	115
333	STAT2001	120
111	BUSN2011	110
111	ECON2102	120
333	BUSN2011	130

UID	CourseNo	Hours
111	COMP2400	120
111	BUSN2011	
111	ECON2102	

STUDY		
UID	CourseNo	Hours
0		120

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

UID	STUDY	
UID	CourseNo	Hours
111	COMP2400	120
222	COMP2400	115
333	STAT2001	120
111	BUSN2011	110
111	ECON2102	120
333	BUSN2011	130



(without any intermediate relation)

STUDY		
UID	CourseNo	Hours
111	COMP2400	120



Rule-based Optimisation

- Can they be executed in one go? ↗ Merging RA operators

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can they be executed in one go? \hookrightarrow Merging RA operators

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can they be executed in one go? \hookrightarrow Merging RA operators

Assignment Project Exam Help

$$\pi_X(\pi_Y(R)) \equiv \pi_X(R) \text{ if } X \subseteq Y,$$

$\pi_{UID}(\pi_{UI}$

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can they be executed in one go? \hookrightarrow Merging RA operators

Assignment Project Exam Help

$$\pi_X(\pi_Y(R)) \equiv \pi_X(R) \text{ if } X \subseteq Y,$$

$\pi_{UID}(\pi_{UI}$

UID		
111		
222	COMP2400	115
333	STAT2001	120
111	BUSN2011	110
111	ECON2101	120
333	BUSN2011	130

333	STAT2001
111	BUSN2
111	ECON2
33	BUSN2

ID
111

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can they be executed in one go? \hookrightarrow Merging RA operators

Assignment Project Exam Help

$$\pi_X(\pi_Y(R)) \equiv \pi_X(R) \text{ if } X \subseteq Y,$$

$\pi_{UID}(\pi_{UI}$

UID	STUDY	
UID	CourseNo	Hours
111	COMP2400	115
333	STAT2001	120
111	BUSN2011	110
111	ECON2102	120
333	BUSN2011	130

333	STAT2001
111	BUSN2
111	ECON2
333	BUSN2

UID
111
222
333

Add WeChat edu_assist_pr

UID	CourseNo	Hours
111	COMP2400	120
222	COMP2400	115
333	STAT2001	120
111	BUSN2011	110
111	ECON2102	120
333	BUSN2011	130



(without any intermediate relation)

UID
111
222
333



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? → Merging RA operators

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? \rightarrow Merging RA operators
 - $\sigma_\varphi(R_1 \sqcap R_2) = R_1 \bowtie_\varphi R_2$

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? \rightarrow Merging RA operators
 - $\sigma_\varphi(R_1 \sqcap R_2) = R_1 \bowtie_\varphi R_2$

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? \rightarrow Merging RA operators

- $\sigma_\varphi(R_1 \quad R_2) \quad R_1 \bowtie_\varphi R_2$

No		
COMP2		
BUSN2011	Management Accounting	6

		Code	Start Date	Status
222	COMP2400	2016 S1	active	active

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? \rightarrow Merging RA operators

$$\bullet \sigma_\varphi(R_1 \quad R_2) \quad R_1 \bowtie_\varphi R_2$$

No	Course	Credit Points
COMP2		
BUSN2011	Management Accounting	6

No	Course	Credit Points	Student ID	Student Name	Unit ID	Unit Name	Start Date	End Date	Status
222	COMP2400	6	111	John Doe	COMP2400	Relational Databases	2016 S1	2016 S2	active
			222	Jane Smith	COMP2400	Relational Databases	2016 S1	2016 S2	active
			111	Samuel Lee	COMP2400	Relational Databases	2016 S2	2017 S1	active

No	Course	Credit Points	Student ID	Student Name	Unit ID	Unit Name	Start Date	End Date	Status
COMP2400	Relational Databases	6	111	John Doe	COMP2400	Relational Databases	2016 S1	2016 S2	active
COMP2400	Relational Databases	6	222	Jane Smith	COMP2400	Relational Databases	2016 S1	2016 S2	active
COMP2400	Relational Databases	6	111	Samuel Lee	COMP2400	Relational Databases	2016 S2	2017 S1	active
BUSN2011	Management Accounting	6	111	John Doe	BUSN2011	Management Accounting	2016 S1	2016 S2	active
BUSN2011	Management Accounting	6	222	Jane Smith	COMP2400	Relational Databases	2016 S1	2016 S2	active
BUSN2011	Management Accounting	6	111	Samuel Lee	COMP2400	Relational Databases	2016 S2	2017 S1	active



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? \rightarrow Merging RA operators

$$\bullet \sigma_\varphi(R_1 \quad R_2) \quad R_1 \bowtie_\varphi R_2$$

No	Course	Credit Points
COMP2		
BUSN2011	Management Accounting	6

No	Course	Credit Points	Student ID	Unit ID	Start Date	Status
222	COMP2400	6	111	COMP2400	2016 S1	active
			222			active
						active

No	Course	Credit Points	Student ID	Unit ID	Start Date	Status
COMP2400	Relational Databases	6	111			e
COMP2400	Relational Databases	6	222			e
COMP2400	Relational Databases	6	111	COMP2400	2016 S2	active
BUSN2011	Management Accounting	6	111	BUSN2011	2016 S1	active
BUSN2011	Management Accounting	6	222	COMP2400	2016 S1	active
BUSN2011	Management Accounting	6	111	COMP2400	2016 S2	active



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? \rightarrow Merging RA operators

- $\sigma_\varphi(R_1 \quad R_2) \quad R_1 \bowtie_\varphi R_2$

No		
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

Course ID	Course Name	Offer Year	Status
222	COMP2400	2016 S1	active
			active
			active

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can they be executed in one go? \rightarrow Merging RA operators

$$\bullet \sigma_\varphi(R_1 \quad R_2) \quad R_1 \bowtie_\varphi R_2$$

<https://eduassistpro.github.io>

No	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

No	Cname	Unit	StudentID	CourseNo	Semester	Status
222	COMP2400	2016 S1				active
						active
						active

Add WeChat `edu_assist_pro`

Inner Join on Course.No=Enrol.CourseNo (no intermediate)

No	Cname	Unit	StudentID	CourseNo	Semester	Status
COMP2400	Relational Databases	6	222	COMP2400	2016 S1	active
COMP2400	Relational Databases	6	111	COMP2400	2016 S2	active
BUSN2011	Management Accounting	6	111	BUSN2011	2016 S1	active



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? ↪ Push select/project before join.

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join.

•

$R_1;$

<https://eduassistpro.github.io>

- $\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}($

attributes not in X , where X_i contain

ones both in R_1 and R_2 ; , and

Add WeChat edu_assist_pro

- $\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , where X_i contains attributes both in R_i and X , and ones both in R_1 and R_2 ;



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

- $\sigma_\varphi(R_1 \bowtie R_2) = \sigma_\varphi(R_1) \bowtie R_2$, if φ contains only attributes in R_1 ;

Course		
COMP2		
BUSN2011	Management Accounting	6

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

- $\sigma_\varphi(R_1 \bowtie R_2) = \sigma_\varphi(R_1) \bowtie R_2$, if φ contains only attributes in R_1 ;

Course		
COMP2		
BUSN2011	Management Accounting	6

Course	Code	Offered	Status
COMP2	COMP2400	2016 S1	active
1			active

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

- $\sigma_\varphi(R_1 \bowtie R_2) = \sigma_\varphi(R_1) \bowtie R_2$, if φ contains only attributes in R_1 ;

Course	Cr	Chi	Status
COMP2			active
BUSN2011	Management Accounting	6	active

Course	Cr	Chi	Status
222	COMP2400	2016 S1	active
1			active

Course	Cr	Chi	Status
COMP2400	Relational Databases	6	active
COMP2400	Relational Databases	6	active
BUSN2011	Management Accounting	6	111



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

- $\sigma_\varphi(R_1 \bowtie R_2) = \sigma_\varphi(R_1) \bowtie R_2$, if φ contains only attributes in R_1 ;

Course	Cname	Unit
COMP2		
BUSN2011	Management Accounting	6

CourseNo	Cname	Unit	StudentID	Semester	Status
222	COMP2400	6		2016 S1	active
1					active

CourseNo	Cname	Unit	StudentID	Semester	Status
COMP2400	Relational Databases	6			
COMP2400	Relational Databases	6			
BUSN2011	Management Accounting	6	111	2016 S1	active

CourseNo	Cname	Unit	StudentID	Semester	Status
BUSN2011	Management Accounting	6	111	2016 S1	active



Rule-based Optimisation

Assignment Project Exam Help

- Can I run the last command? ↪ Push select/project before you run

-

$R_1;$

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can I run the project last? ↪ Push select/project before icon

-

$R_1;$

Course		
COMP2		
BUSN2011	Management Accounting	6

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can I run the project last? ↪ Push select/project before icon

-

$R_1;$

<https://eduassistpro.github.io>

Course		
COMP2		
BUSN2011	Management Accounting	6

COURSE		
CourseNo	Cname	Unit
BUSN2011	Management Accounting	6

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

-

$R_1;$

<https://eduassistpro.github.io>

Course		
COMP2		
BUSN2011	Management Accounting	6

COURSE		
CourseNo	CourseName	Unit
BUSN2011	Management Accounting	6

StudentID	CourseNo	Semester	Status
			active
			active
			active

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

$R_1;$

<https://eduassistpro.github.io>

Course		
COMP2		
BUSN2011	Management Accounting	6

COURSE		
CourseNo	Cname	Unit
BUSN2011	Management Accounting	6

StudentID	CourseNo	Semester	Status
			active
			active
			active

Add WeChat edu_assist_pro

CourseNo	Cname	Unit	StudentID	Semester	Status
BUSN2011	Management Accounting	6	111	2016 S1	active



Rule-based Optimisation

- Can join be executed last? ↗ Push select/project before join.

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

Assignment Project Exam Help

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

Assignment Project Exam Help

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

Course		
COMP2		
BUSN2011	Management Accounting	6

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

Course		
COMP2		
BUSN2011	Management Accounting	6

		Ter	Stats
222	COMP2400	1	active
		2016 S1	active
			active

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Course		
COMP2		
BUSN2011	Management Accounting	6

			Ter	Stats
			1	active
222	COMP2400	2016 S1		active
				active

Course No	Course Name	Unit			
COMP2400	Relational Databases	6			
COMP2400	Relational Databases	6			
BUSN2011	Management Accounting	6	111	2016 S1	active



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Course		
COMP2		
BUSN2011	Management Accounting	6

			Ter	Stats
			1	active
222	COMP2400	2016 S1		active
				active

CourseNo	Cname	Unit		
COMP2400	Relational Databases	6		
COMP2400	Relational Databases	6		
BUSN2011	Management Accounting	6	111	2016 S1 active

CourseNo	Cname	StudentID
COMP2400	Relational Databases	222
COMP2400	Relational Databases	111
BUSN2011	Management Accounting	111



Rule-based Optimisation

- Can join be executed last? ↗ Push select/project before join.

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

Assignment Project Exam Help

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only

attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

Assignment Project Exam Help

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only

attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Course		
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Course		COURSE
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

CourseNo	CourseName	COURSE
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Course	COURSE	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

CourseNo	CourseName	Unit	Status
222	COMP2400	2016 S1	active
			active

CourseNo	CourseName	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Course	COURSE	Stu	Status
COMP2400	Relational Databases	6	active
BUSN2011	Management Accounting	6	active

CourseNo	CourseName	Stu	Status
222	COMP2400	2016 S1	active
			active

CourseNo	CourseName	Stu	Status
COMP2400	Relational Databases	6	active
BUSN2011	Management Accounting	6	active

CourseNo	CourseName	Stu	Status
111	COMP2400	2016 S1	active
			active

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)$, if the join condition involves only attributes in X , how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Course	Cname	StudentID
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

CourseNo	Cname	StudentID	Status
222	COMP2400	2016 S1	active
			active
			active

CourseNo	Cname
COMP2400	Relational Databases
BUSN2011	Management Accounting

CourseNo	Cname
222	COMP2400
111	COMP2400

CourseNo	Cname	StudentID
COMP2400	Relational Databases	222
COMP2400	Relational Databases	111
BUSN2011	Management Accounting	111



Rule-based Optimisation

- Can join be executed last? ↗ Push select/project before join.

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

Assignment Project Exam Help

$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

Assignment Project Exam Help

$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

Assignment Project Exam Help

$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

Course		
COMP2		
BUSN2011	Management Accounting	6

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

Course		
COMP2		
BUSN2011	Management Accounting	6

		Ter	Stats
222	COMP2400	1	active
		2016 S1	active
			active

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Course		
COMP2		
BUSN2011	Management Accounting	6

			Ter	Stats
			1	active
222	COMP2400	2016 S1		active
				active

Course No	Course Name	Unit		
COMP2400	Relational Databases	6		
COMP2400	Relational Databases	6		
BUSN2011	Management Accounting	6	111	2016 S1 active



Rule-based Optimisation

- Can join be executed last? \rightarrow Push select/project before join.

$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io>

Course		
COMP2		
BUSN2011	Management Accounting	6

			Ter	Stats
			1	active
222	COMP2400	2016 S1		active
				active

Course No	Cname	Unit		
COMP2400	Relational Databases	6		
COMP2400	Relational Databases	6		
BUSN2011	Management Accounting	6	111	2016 S1 active

Cname	StudentID
Relational Databases	222
Relational Databases	111
Management Accounting	111

Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

• es

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? ↪ Push select/project before join

-

es

<https://eduassistpro.github.io>

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

-

es

<https://eduassistpro.github.io>

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

π_{Cname} COURSE
Cname
Relational
Management

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

es

$\pi_{Cname} \text{ COURSE}$ \hookrightarrow $\text{tp}(Eprol) ?$

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

StudentID	CourseNo	Semester	Status
			active
			active
			active

$\pi_{Cname} \text{ COURSE}$
Cname
Relational
Management

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

-

es

$\pi_{Cname} \text{ COURSE}$ \hookrightarrow $\text{tp}(Eprol) ?$

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

$\pi_{Cname} \text{ COURSE}$
Cname
Relational
Management

StudentID	CourseNo	Semester	Status
			active
			active
			active

π_{St}
111
222

Add WeChat edu_assist_pro



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

-

es

$\pi_{Cname}(\text{Course}) \bowtie \pi_{StudentID}(\text{Enrol})$?

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

StudentID	CourseNo	Semester	Status
			active
			active
			active

$\pi_{Cname}(\text{Course})$	Cname
	Relational
	Management

$\pi_{StudentID}(\text{Enrol})$	St
	111
	222

Is $\pi_{Cname}(\text{Course}) \bowtie \pi_{StudentID}(\text{Enrol})$ our desired result?



Rule-based Optimisation

Assignment Project Exam Help

- Can join be executed last? \hookrightarrow Push select/project before join

-

es

$\pi_{Cname}(Course) \bowtie \pi_{StudentID}(Enrol)$?

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

StudentID	CourseNo	Semester	Status
			active
			active
			active

$\pi_{Cname}(Course)$	\bowtie	$\pi_{StudentID}(Enrol)$
Cname		
Relational		
Management		

Add WeChat edu_assist_pro

Is $\pi_{Cname}(Course) \bowtie \pi_{StudentID}(Enrol)$ our desired result?

No. $\pi_{Cname, StudentID}(Course \bowtie Enrol) \neq \pi_{Cname}(Course) \bowtie \pi_{StudentID}(Enrol)$



Rule-based Optimisation

- Can join be executed last? \hookrightarrow Push select/project before join.

Assignment Project Exam Help

$\pi_X(A_1 \bowtie A_2) = \pi_X(\pi_{X_1}(A_1) \bowtie \pi_{X_2}(A_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?
nrol()

Add WeChat edu_assist_pr



Rule-based Optimisation

- Can join be executed last? \hookrightarrow Push select/project before join.

Assignment Project Exam Help

$\pi_{X \cup Y}(\rho_1 \bowtie \rho_2) = \pi_X(\pi_{X_1}(\rho_1) \bowtie \pi_{X_2}(\rho_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io/nrol/>

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \hookrightarrow Push select/project before join.

$\pi_X(P_1 \bowtie P_2) = \pi_X(\pi_{X_1}(P_1) \bowtie \pi_{X_2}(P_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

nrol()

<https://eduassistpro.github.io>

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

$\pi_{\{CourseNo, Cname\}}(COURSES)$	
CourseNo	Cname
COMP2400	Relational Databases
BUSN2011	Management Accounting

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \hookrightarrow Push select/project before join.

$\pi_{X \cup Y}(\rho_{X_1 \times X_2}(A_1 \times A_2)) = \pi_X(\pi_{X_1}(A_1) \times \pi_{X_2}(A_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

$\pi_{\{CourseNo, Cname\}}(COURSES)$	
CourseNo	Cname
COMP2400	Relational Databases
BUSN2011	Management Accounting

		ter	Status
111	BUSN2011	2016 S1	active
			active
			active

Add WeChat edu_assist_pro



Rule-based Optimisation

- Can join be executed last? \hookrightarrow Push select/project before join.

$\pi_{X \cup Y}(\pi_X(A_1) \bowtie \pi_Y(A_2)) = \pi_X(\pi_{X_1}(A_1) \bowtie \pi_{X_2}(A_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

$\pi_{\{CourseNo, Cname\}}(COURSES)$	
CourseNo	Cname
COMP2400	Relational Databases
BUSN2011	Management Accounting

nro()	ter	Status
111	BUSN2011	2016 S1
		active
		active
		active

$\pi_{\{CourseNo\}}(Stu)$	
222	COMP2400
111	COMP2400

Add WeChat [edu_assist_pro](https://eduassistpro.github.io)



Rule-based Optimisation

- Can join be executed last? \hookrightarrow Push select/project before join.

$\pi_{X \cup Y}(\pi_X(A_1) \bowtie \pi_Y(A_2)) = \pi_X(\pi_{X_1}(A_1) \bowtie \pi_{X_2}(A_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io/nrof/>

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

$\pi_{\{CourseNo, Cname\}}(COURSES)$	
CourseNo	Cname
COMP2400	Relational Databases
BUSN2011	Management Accounting

CourseNo	StuID	Year	Status
111	BUSN2011	2016 S1	active
			active
			active

$\pi_{\{CourseNo, StuID\}}(STUDENTS)$	
222	COMP2400
111	COMP2400

CourseNo	Cname	StudentID
COMP2400	Relational Databases	222
COMP2400	Relational Databases	111
BUSN2011	Management Accounting	111



Rule-based Optimisation

- Can join be executed last? \hookrightarrow Push select/project before join.

$\pi_{X \cup Y}(\pi_X(A_1) \bowtie \pi_Y(A_2)) = \pi_X(\pi_{X_1}(A_1) \bowtie \pi_{X_2}(A_2))$, if the join condition involves attributes outside X, how could we derive X_1 and X_2 ?

<https://eduassistpro.github.io/nrof/>

CourseNo	Cname	Unit
COMP2400	Relational Databases	6
BUSN2011	Management Accounting	6

CourseNo	Cname
COMP2400	Relational Databases
BUSN2011	Management Accounting

CourseNo	Cname	StuID	Year	Status
111	BUSN2011		2016 S1	active
				active
				active

CourseNo	Cname
222	COMP2400
111	COMP2400

CourseNo	Cname	StudentID
COMP2400	Relational Databases	222
COMP2400	Relational Databases	111
BUSN2011	Management Accounting	111

Cname	StudentID
Relational Databases	222
Relational Databases	111
Management Accounting	111



Heuristic Rules and Query Trees

(1) $\sigma_\varphi(\sigma_\psi(R)) \equiv \sigma_{\varphi \wedge \psi}(R)$

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Heuristic Rules and Query Trees

$$(1) \sigma_{\varphi}(\sigma_{\psi}(R)) \equiv \sigma_{\varphi \wedge \psi}(R)$$

Assignment Project Exam Help

<https://eduassistpro.github.io>

$$(2) \pi_X(\pi_Y(R)) \equiv \pi_X(R) \text{ if } X \subseteq Y;$$

Add WeChat edu_assist_pro



Heuristic Rules

$$(3) \sigma_\varphi(R_1 \times R_2) \equiv R_1 \bowtie_\varphi R_2$$

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Heuristic Rules

$$(3) \sigma_\varphi(R_1 \times R_2) \equiv R_1 \bowtie_\varphi R_2$$

Assignment Project Exam Help

<https://eduassistpro.github.io>

$$(4) \sigma_{\varphi_1}(R_1 \bowtie_{\varphi_2} R_2) \equiv R_2 \bowtie_{\varphi_1 \wedge \varphi_2} R_1$$

Add WeChat edu_assist_pro



Push-down Selection – Example

Assignment Project Exam Help

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

MOVIE(title, production_year, country, run_time, major_genre)

R

- $\pi_{tit} \circ \text{https://eduassistpro.github.io}$

π_{tit}

$first_name = 'Tom' \wedge last_name = 'Cruise'$ (MOVIE \times (PE

Add WeChat edu_assist_pro



Push-down Selection – Example

Assignment Project Exam Help

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

MOVIE(title, production_year, country, run_time, major_genre)

R

- π_{tit}^q <https://eduassistpro.github.io>

π_{tit}

$first_name = 'Tom' \wedge last_name = 'Cruise'$ (MOVIE \times (PE

- Question: Can we apply the following rule to optimise

$\sigma_{\varphi_1 \wedge \varphi_2}(R_1 \times R_2) \equiv \sigma_{\varphi_1}(R_1) \times \sigma_{\varphi_2}(R_2)$ if

and φ_2 contains only attributes in R_2

Add WeChat edu_assist_pro



Push-down Selection – Example

Assignment Project Exam Help

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

MOVIE(title, production_year, country, run_time, major_genre)

R

- $\pi_{tit}^q \text{https://eduassistpro.github.io}$

$\pi_{tit}^q \text{first_name='Tom' \wedge last_name='Cruise'} (\text{MOVIE} \times (\text{PE}$

- Question: Can we apply the following rule to optimize?

$\sigma_{\varphi_1 \wedge \varphi_2}(R_1 \times R_2) \equiv \sigma_{\varphi_1}(R_1) \times \sigma_{\varphi_2}(R_2)$ if R_1

and φ_2 contains only attributes in R_2

- We would have

$\pi_{title, production_year}(\sigma_{title=mtitle \wedge production_year=mprod_year}(\sigma_{major_genre='war'}(\text{MOVIE}))$

$\times \sigma_{first_name='Tom' \wedge last_name='Cruise'}(\text{PERSON} \bowtie \text{ROLE}))$



Push-down Selection – Example

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Push-down Selection – Example

Assignment Project Exam Help

- # Assignment Project

MOVIE

B

- <https://eduassistpro.github.io>

$\pi_{title, production_year}(\sigma_{title=mtitle \wedge production_year=mpy} - (\sigma_{$

$\times \sigma_{first_name} = 'Tom' \wedge \sigma_{last_name} = 'Cruise'$ (PERSON ▷

〈*σ first_name = 'Tom' ∧ last_name = 'Cruise'* ▷ PERSON▷



Push-down Selection – Example

Assignment Project Exam Help

- Given the relation schema:
 $\text{PERSON}(\text{id}, \text{first_name}, \text{last_name}, \text{year_born})$

MOVIE

R

- $\sigma_{\varphi} \text{https://eduassistpro.github.io}$

$\pi_{\text{title, production_year}}(\sigma_{\text{title} = m\text{title} \wedge \text{production_year} = m\text{pr}} \text{MOVIE})$

$\times \sigma_{\text{first_name} = 'Tom' \wedge \text{last_name} = 'Cruise'} \text{PERSON} \triangleright$

- Can we apply $\sigma_{\varphi}(R_1 \times R_2) \sqsubseteq R_1 \bowtie_{\varphi} R_2$?

Add WeChat edu_assist_pro



Push-down Selection – Example

Assignment Project Exam Help

- Given the relation schema:
 $\text{PERSON}(\text{id}, \text{first_name}, \text{last_name}, \text{year_born})$

MOVIE

R

- $\sigma_{\varphi} \text{https://eduassistpro.github.io}$

$$\pi_{\text{title, production_year}}(\sigma_{\text{title} = m\text{title} \wedge \text{production_year} = m\text{prod_year}} \text{MOVIE})$$

$$\times \sigma_{\text{first_name} = 'Tom' \wedge \text{last_name} = 'Cruise'}(\text{PERSON} \bowtie \text{ROLE})$$

- Can we apply $\sigma_{\varphi}(R_1 \times R_2) \equiv R_1 \bowtie_{\varphi} R_2$?
- We would have

$$\pi_{\text{title, production_year}}(\sigma_{\text{major_genre} = 'war'}(\text{MOVIE}) \bowtie_{\text{title} = m\text{title} \wedge \text{production_year} = m\text{prod_year}} (\text{PERSON} \bowtie \text{ROLE}))$$

$$\sigma_{\text{first_name} = 'Tom' \wedge \text{last_name} = 'Cruise'}(\text{PERSON} \bowtie \text{ROLE}))$$

Add WeChat edu_assist_pro



Push-down Selection – Example

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Push-down Projection – Example

Assignment Project Exam Help

- Given the relation schema:

PERSON(id, first_name, last_name, year_born)

M

R

Q <https://eduassistpro.github.io>

$\pi_{title, production_year}(\sigma_{major_genre='war'}(MOVIE) \bowtie$

$\sigma_{first_name='Tim' \wedge last_name='Pruett'}(PERSON) \bowtie$

Add WeChat edu_assist_pro



Push-down Projection – Example

Assignment Project Exam Help

- Given the relation schema:

PERSON(id, first_name, last_name, year_born)

M

R

Q <https://eduassistpro.github.io>

$\pi_{title, production_year}(\sigma_{major_genre='war'}(MOVIE) \bowtie$

$\sigma_{first_name='Tim' \wedge last_name='Bruland'}(PERSON) \bowtie$

- Question: Can we apply the following rule to optimi

$$\pi_X(R_1 \bowtie R_2) \equiv \pi_X(\pi_{X_1}(R_1) \bowtie \pi_{X_2}(R_2)),$$

where X_i contains attributes both in R_i and X , and ones both in R_1 and R_2



Push-down Projection – Example

Assignment Project Exam Help

- Given the relation schemas:

PERSON(id, first_name, last_name, year_born)

MOVIE(title, production_year, country, run_time, major_genre)

R

- σ <https://eduassistpro.github.io>

$\pi_{title, production_year}(\sigma_{major_genre='war'}(\text{MOVIE}) \bowtie_{title=mtitle \ production_year=mprod_year} ($

$\sigma_{first_name='Tom' \wedge last_name='Cruise'}(\text{PERSON} \bowtie$

Add WeChat edu_assist_pro



Push-down Projection – Example

Assignment Project Exam Help

- Given the relation schema:

PERSON(id, first_name, last_name, year_born)

MOVIE(title, production_year, country, run_time, major_genre)

R

- <https://eduassistpro.github.io>

$\pi_{title, production_year}(\sigma_{major_genre='war'}(\text{MOVIE}) \bowtie_{title=mtitle \ production_year=mprod_year} ($

$\sigma_{first_name='Tom' \wedge last_name='Cruise'}(\text{PERSON} \bowtie$

- We would have:

$\pi_{title, production_year}(\pi_{title, production_year}(\sigma_{major_genre='war'}$

$\bowtie_{title=mtitle \wedge production_year=mprod_year}$

$(\pi_{mtitle, mprod_year}(\sigma_{first_name='Tom' \wedge last_name='Cruise'}(\text{PERSON} \bowtie \text{ROLE}))))$

We further apply some rules to optimise the query ...



Push-down Projection – Example

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro



Cost-based Optimisation (not assessed)

- Consider CHARTS = {Rank, Artist, Song} with 100 tuples and 3 attributes

Rank	Artist	Song
1	Chingy	Right Thurr
6	Brooke Fraser	
7	Black Eyed Peas	
...	...	
...	...	

- Compare two strategies of evaluating “Who is top of the pops?”:

- $\sigma_{\text{Rank}=1}(\pi_{\text{Rank}, \text{Artist}}(\text{CHARTS}))$
- $\pi_{\text{Rank}, \text{Artist}}(\sigma_{\text{Rank}=1}(\text{CHARTS}))$



Cost-based Optimisation (not assessed)

- Consider CHARTS = {Rank, Artist, Song} with 100 tuples and 3 attributes

Rank	Artist	Song
1	Chingy	Right Thurr
6	Brooke Fraser	
7	Black Eyed Peas	
...	...	
...	...	

- Compare two strategies of evaluating “Who is top of the pops?”:

- $\sigma_{\text{Rank}=1}(\pi_{\text{Rank}, \text{Artist}}(\text{CHARTS}))$
- $\pi_{\text{Rank}, \text{Artist}}(\sigma_{\text{Rank}=1}(\text{CHARTS}))$

Selection before Projection is preferred.



Cost-based Optimisation (not assessed)

- Consider CHARTS = {Rank, Artist, ...} with 100 tuples and 50 attributes:

Rank	Artist	Song
			
			
			
5	Justin Timberlake	Senorita
6	Brooke Fraser				
7	Black Eyed Peas	Where Is The Love			
...			

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

- Compare two strategies of evaluating?
- $\sigma_{\text{Rank} > 10}(\pi_{\text{Rank}, \text{Artist}}(\text{CHARTS}))$
- $\pi_{\text{Rank}, \text{Artist}}(\sigma_{\text{Rank} > 10}(\text{CHARTS}))$



Cost-based Optimisation (not assessed)

- Consider CHARTS = {Rank, Artist, ...} with 100 tuples and 50 attributes:

Rank	Artist	Song
			
			
			
5	Justin Timberlake	Senorita
6	Brooke Fraser				
7	Black Eyed Peas	Where Is The Love			
...			

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

- Compare two strategies of evaluating?

 - $\sigma_{\text{Rank} > 10}(\pi_{\text{Rank}, \text{Artist}}(\text{CHARTS}))$
 - $\pi_{\text{Rank}, \text{Artist}}(\sigma_{\text{Rank} > 10}(\text{CHARTS}))$

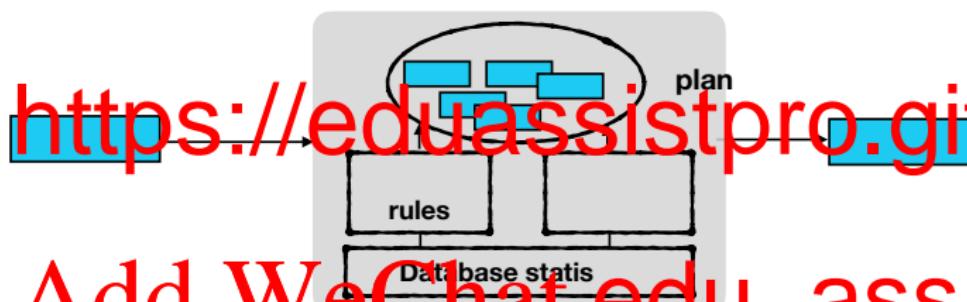
Projection before Selection is preferred.



Query Optimisation

Assignment Project Exam Help

Query Optimiser



<https://eduassistpro.github.io>

Add WeChat `edu_assist_pro`

- Trade-off:

Time for executing a RA query vs Time for finding a better RA query



(credit cookie) memorising vs understanding

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro