



Week 10 Workshop - Database Transactions

# Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro



## Transactions

- A **transaction** is a sequence of database operations grouped together for execution as a logic unit in a DBMS.

# Assignment Project Exam Help

<https://eduassistpro.github.io>

5

COMMIT;

Add WeChat edu\_assist\_pro



## Transactions

- What's the difference between database transactions and programs written by a programming language like C, Java, and Python?
- How are transactions handled in the query processing?

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro

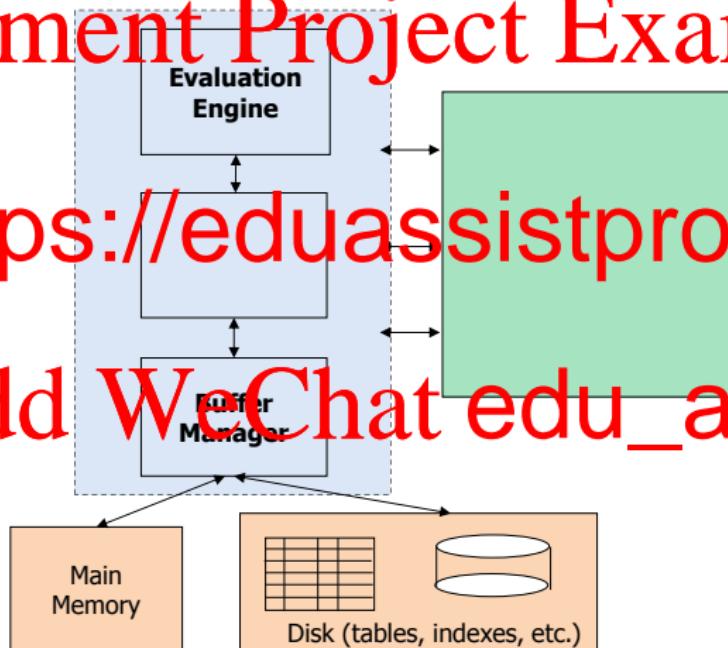


## Transaction Manager - A Simplified View

# Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro





## Transactions - ACID Properties

# Assignment Project Exam Help

Transactions

ACID properties

<https://eduassistpro.github.io>

$T_2$  : SELECT ...

$T_3$  : INSERT ...

$T_4$  : BEGIN TRANSACTION

SELECT ...

DELETE ...

ABORT

Consistency

Add WeChat edu\_assist\_pr

Durability



## Transactions - ACID Properties

# Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro



## Transactions - ACID Properties

# Assignment Project Exam Help

ACID properties

Transaction Manager

<https://eduassistpro.github.io>

Isolation

Durability

Add WeChat edu\_assist\_pro

Consistency is the responsibility of an application developer.



## Transaction Manager - Common Techniques

# Assignment Project Exam Help

Transaction Manager

<https://eduassistpro.github.io>

- Add WeChat edu\_assist\_pr  
e.g., Write-Ahead Log (WAL) Protocol



## Logging - Introduction

# Assignment Project Exam Help

- A transaction log is an append-only file that records changes to objects made by transactions.

- W
- A t  
da

<https://eduassistpro.github.io>

- Recovery amounts to either undoing or red

o Undo the operations that have not been co

- Redo the operations that have been com  
written to disk.

- Checkpoints tell the points from which to begin applying transaction logs during database recovery.



## Write-Ahead Log (WAL) Protocol

# Assignment Project Exam Help

- Write-Ahead Log (WAL) requires that a record of every change to a database is available while attempting to recover from a crash.

<https://eduassistpro.github.io>

- A record in the log must be written to persistent before committing the transaction.
- Accordingly, the definition of a **commit** record is:

**“A transaction, all of whose log records, including a commit record, have been written to persistent storage”.**

Add WeChat [edu\\_assist\\_pro](https://edu_assist_pro)



## Write-Ahead Log (WAL) Protocol

- Typical fields in a log record:



- Each log record has a unique id called **LSN** (Log

- Possible **types** include: update, commit, abort, end, etc.

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pr

- Does WAL bring in some benefits for performance

- Often results in a significantly reduced number of disk writes
- Supports one sync against the log file instead of potentially many against the data files
- Enables online backup and point-in-time recovery



## Transaction Manager - Recovery

# Assignment Project Exam Help

- Key concepts to aid in recovery:

• <https://eduassistpro.github.io/>

- Undo ...

- Redo ...

Add WeChat edu\_assist\_pro

- **Checkpoint:** snapshot of the state of a data

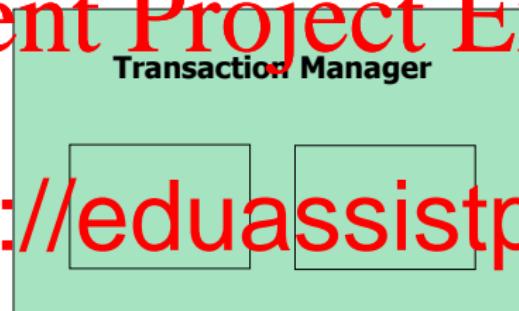
(Widely used in practice, but not covered in this course)



## Transaction Manager - Common Techniques

# Assignment Project Exam Help

<https://eduassistpro.github.io>



- **Add WeChat** [edu\\_assist\\_pro](https://edu_assist_pro)
  - e.g., Write-Ahead Log (WAL) Protocol
- **Locking** for concurrency control – assuring **isolation** of transactions
  - e.g., Two-Phase Locking (2PL) Protocol



## Locking - Introduction

# Assignment Project Exam Help

- A lock is associated with an object, e.g., file, table, record, page, etc.

- T

<https://eduassistpro.github.io>

(Note: there are other types of locks defined by differ

- Add WeChat edu\_assist\_pr

Lock type	read-lock	write-lock
read-lock	Yes	No
write-lock	No	No



## Two-Phase Locking (2PL) Protocol

# Assignment Project Exam Help

- Locks are handled in two phases:

- 
- 

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro

## Two-Phase Locking (2PL) Protocol

# Assignment Project Exam Help

Bad news:

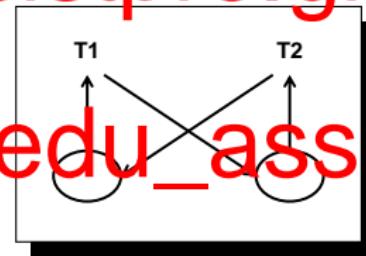
- 2PL can radically limit interleaving among transactions in some cases ...

- 2

tra

<https://eduassistpro.github.io>

	1	2
1	lock-r(A)	
2	read(A)	
3		lock-r(B) read(B)
4		
5	lock-w(B)	
6	write(B)	
7		lock-w(A) write(A)
8		



- $T_1$  is waiting for  $T_2$  to get a write-lock on  $B$ .  $T_2$  is waiting for  $T_1$  to get a write-lock on  $A$ .



## Two-Phase Locking (2PL) Protocol

# Assignment Project Exam Help

Good

- 2 tra <https://eduassistpro.github.io>
  - **Serializability** means that a resulting database state of running transactions ser
  - Serializability is the major concern in transactions.

Add WeChat [edu\\_assist\\_pro](https://edu_assist_pro)



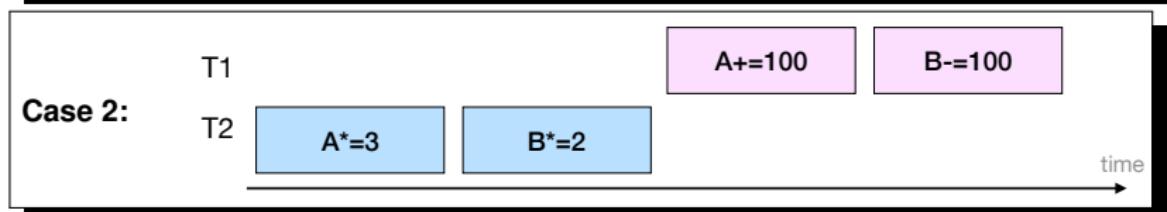
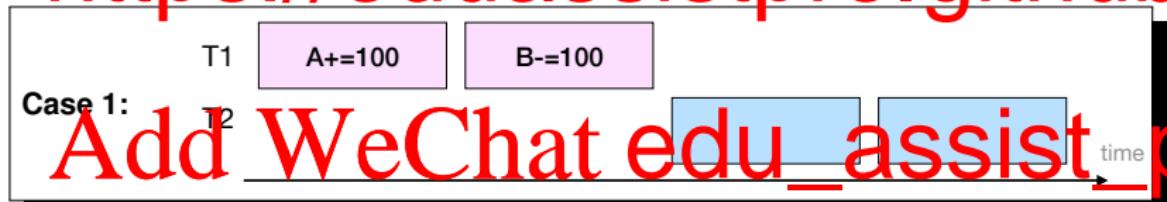
## Serializability - Example

- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions.

# Assignment Project Exam Help



- S <https://eduassistpro.github.io>

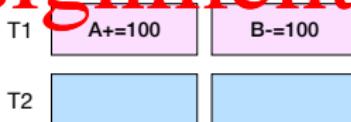




## Serializability - Example

- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions.

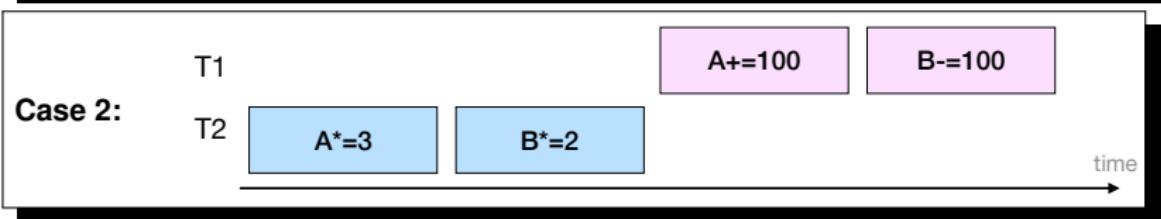
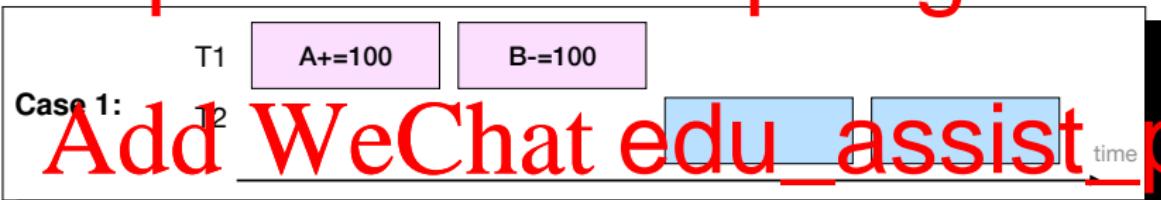
# Assignment Project Exam Help



- Case 1:  $A=900$  and  $B=800$

•

- S <https://eduassistpro.github.io>

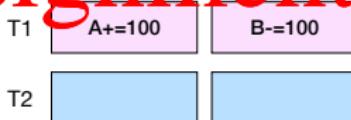




## Serializability - Example

- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions.

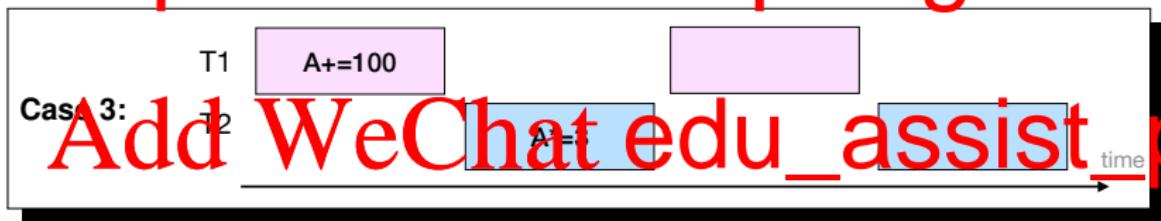
# Assignment Project Exam Help



- Case 1:  $A=900$  and  $B=800$

$\vdots$

- Another <https://eduassistpro.github.io>



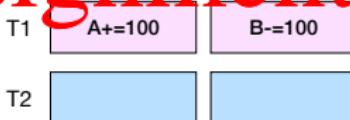
- Yes.  $A=900$  and  $B=800 \leftrightarrow$  equivalent to Case 1!



## Serializability - Example

- Consider  $A = 200$  and  $B = 500$ , and we have two concurrent transactions.

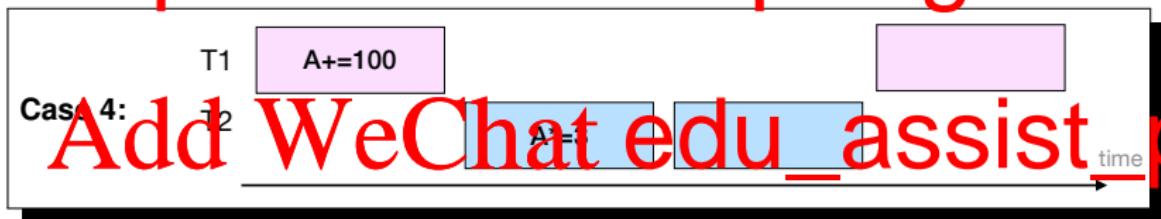
# Assignment Project Exam Help



- Case 1:  $A=900$  and  $B=800$

$\vdots$

- Another <https://eduassistpro.github.io>



- No.  $A=900$  and  $B=900 \rightarrow$  not equivalent to Case 1 or Case 2!



## Problems in Concurrent Transactions

# Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:

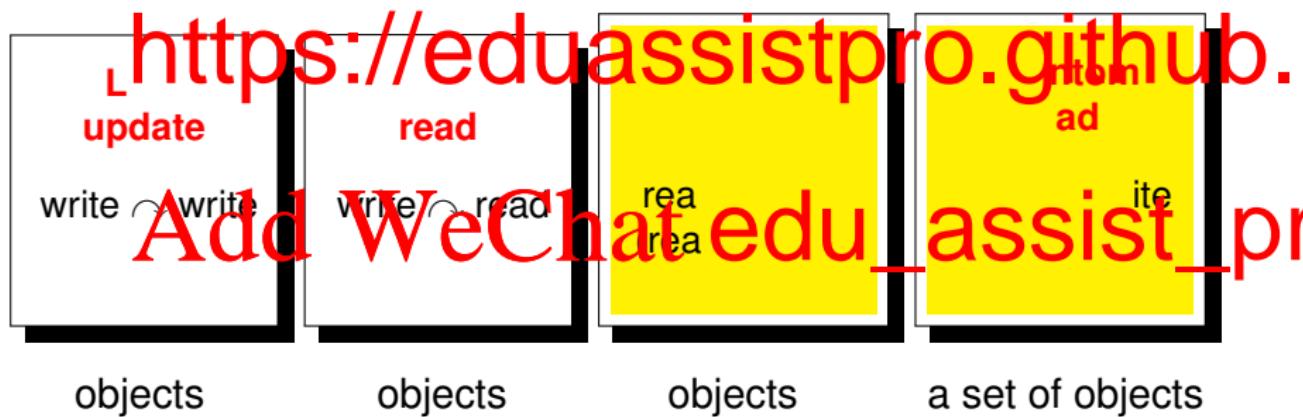




## Problems in Concurrent Transactions

# Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:





## The Lost Update Problem - Another Example

- Ben and Amy have the same salary.  $T_1$  sets their salaries to \$80,000, and  $T_2$  sets their salaries to \$90,000.

# Assignment Project Exam Help

- ① If executing  $T_1$  and  $T_2$  sequentially,

• <https://eduassistpro.github.io>

- ② If executing  $T_1$  and  $T_2$  concurrently,

• Add WeChat edu\_assist\_pr

	$\bar{T}_1$	$\bar{T}_2$
1	write(A) (A:=80000)	
2		write(A) (A:=
3		write(B) (B:=90000)
4		commit
5	write(B) (B:=80000)	
6	commit	

→ It is not acceptable!



## The Dirty Read Problem - Another Example

- Both Ben and Amy are rewarded a bonus \$5,000 and a pay rise 5%.  $T_1$  increases their salaries with \$1,000 and  $T_2$  increments their salaries by 5%

# Assignment Project Exam Help

- If executing  $T_1$  and  $T_2$  sequentially, they would have .... Also,  $T_1$  or  $T_2$

<https://eduassistpro.github.io>

	$T_1$	$T_2$
1	read(A)	
2	write(A) ( $A := A + 500$ )	read(A)
3		
4	read(B)	
5	write(B) ( $B := B + 5000$ )	
6	abort	
7		write(A) ( $A := A + A \times 5\%$ )
8		read(B)
9		write(B) ( $B := B + B \times 5\%$ )
10		commit

→ It is not acceptable!

Add WeChat edu\_assist\_pro



## The Unrepeatable Read Problem - Another Example

# Assignment Project Exam Help

- Amy and Ben are using a website to book flight tickets to Brisbane.

• <https://eduassistpro.github.io>

Add WeChat edu\_assist\_pr

	$T_1$ (from Amy)	$T_2$
1	read(X)	
2		re
3		wr
4		commit
5	read(X)	

- This situation can never arise in a serial execution of  $T_1$  and  $T_2$ .



## The Phantom Read Problem - Another Example

# Assignment Project Exam Help

- Amy is 30 years old, but her age in the table players is mistakenly recorded as 40. Ben is 28 years old and his age is correctly recorded in players.
- Suppose that we have the following two current transactions:

$T$

<https://eduassistpro.github.io>

WHERE age < 32;  
COMMIT;

COMMIT;

Add WeChat [edu\\_assist\\_pro](https://edu_assist_pro)

	$T_1$	$T_2$
1	read(players)	
2		re
3		write(players)
4		commit
5	read(players)	
6	commit	

- This situation also can never arise in a serial execution of  $T_1$  and  $T_2$ .



## Discussion

# Assignment Project Exam Help

- Unrepeatable read

<https://eduassistpro.github.io>

- Can be prevented using record-level locking.

- Phantom read

Add WeChat edu\_assist\_pro

- Executing the same SELECT twice on different tuples;

- May occur when querying a set of tuples that are affected by **INSERT/DELETE/UPDATE** from another transaction;
- Can be prevented using table-level locking.



## What Should We lock?

# Assignment Project Exam Help

- Consider the following two concurrent transactions again:

$T_1: \text{SELECT } * \text{ FROM players}$   
 $\text{WHERE age} < 32;$

$T_2: \text{UPDATE players}$

<https://eduassistpro.github.io>

- What objects should the DBMS lock in order to avoid the problem?

Add WeChat edu\_assist\_pro

- Table-level locks**

e.g., read-lock on players for  $T_1$ , write-lock on

2

- Record-level locks**

e.g., read-lock on every record with  $\text{age} < 32$  for  $T_1$ , write-lock on every record with  $\text{rating} = 8$  and  $\text{name} = \text{'Amy'}$  for  $T_2$

- ...



## Transaction Support in SQL

# Assignment Project Exam Help

- An explicit transaction may have no `BEGIN TRANSACTION` statement, but must be ended with either `COMMIT` or `ABORT (ROLLBACK)` statement.

- W  
st
- T  
th

<https://eduassistpro.github.io>

transaction is prepared to tolerate on concurrent tr

• Key idea:  
**Add WeChat edu\_assist\_pr**

To trade off **consistency** (i.e., increased risk of violating database integrity) with **performance** (i.e., greater concurrent access to data)



## Isolation Levels

# Assignment Project Exam Help

- SQL:92 defines four isolation levels:

1 Read Uncommitted



4 Serializable

- To specify an isolation level, e.g.,

```
SET TRANSACTION ISOLATION LEVEL serializable
```

- The SQL standard does not impose a specific locking scheme or mandate particular behaviors.

# Add WeChat edu\_assist\_pr



## Isolation Levels

# Assignment Project Exam Help

I	R	E	ead
RE			
R			
RE			
SERIALIZABLE			

<https://eduassistpro.github.io>

- Different DBMSs implement isolation levels slightly differently.
- Add WeChat [edu\\_assist\\_pro](https://edu_assist_pro)
- The isolation level required for **Lost Update** is debatable (depending on a DBMS's implementations). But in general, it may require the highest level **SERIALIZABLE** to prevent it.<sup>1</sup>

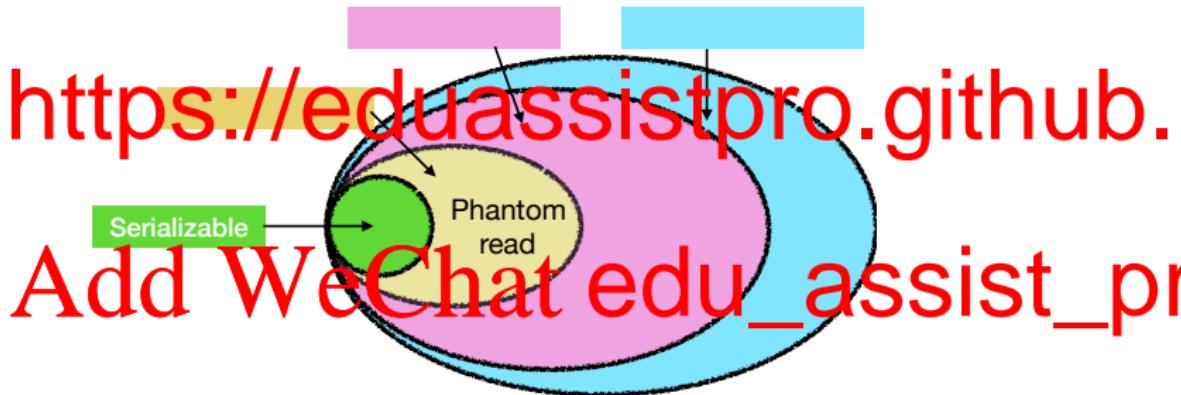
<sup>1</sup> [https://drtom.ch/posts/2011/11/12/The\\_Lost\\_Update\\_Problem\\_-\\_-Part\\_1/](https://drtom.ch/posts/2011/11/12/The_Lost_Update_Problem_-_-Part_1/)



## Isolation Levels - Concurrency Control

# Assignment Project Exam Help

- A DBMS provides different levels of isolation → different degrees of concurrency control to prevent different problems.



- Concurrency control is **NOT** binary in a database system.



## Isolation Levels - Read Uncommitted

# Assignment Project Exam Help

• **Read Uncommitted** is the least restrictive isolation level.

- One transaction can see changes made by other transactions which are not yet committed.

- U  
w

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro



## Isolation Levels - Read Committed

- **Read Committed**: One transaction only sees committed changes by other transactions.

- It is the most commonly used isolation level in database applications.
- U

no <https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro



## Isolation Levels - Repeatable Reads

# Assignment Project Exam Help

- **Repeatable Reads:** The objects touched by a transaction are locked and cannot be updated or deleted by a concurrent transaction.

• U

do

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro



## Isolation Levels - Serializable

# Assignment Project Exam Help

- Serializable is the highest isolation level. All transactions are totally isolated from other transactions. It is safe but may cause significant performance hit.

- U  
do  
dif

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro



Locks Taken by SQL Server for Isolation Levels<sup>2</sup>

# Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro

<sup>2</sup><http://michaeljswart.com/2012/06/visualizing-transaction-isolations-for-sql-server/>



## Wrap-up - Isolation Levels

# Assignment Project Exam Help

- A lower isolation level increases the ability of many users to access data at the same time, but also increases the number of concurrency effects (s

- Cef https://eduassistpro.github.io and increases the chances that one transaction will block another.

- Choosing the appropriate isolation level depends

Add WeChat edu\_assist\_pro

- the data integrity requirements of the application

against

- the overhead of each isolation level.



# Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro



## Research Topics

# Assignment Project Exam Help

- This is an active research area covering many interesting research topics.

- Hi
- da
- H

<https://eduassistpro.github.io>

system is a relational database system or somethi

Add WeChat edu\_assist\_pr

- Distributed database systems
- Graph database systems
- Document-oriented database systems
- ...



## Research Topics

Assignment Project Exam Help

- Distributed transactions

<https://eduassistpro.github.io>

Add WeChat edu\_assist\_pro